

Working with Caching and Concurrency



Kevin Dockx

ARCHITECT

@KevinDockx <https://www.kevindockx.com>



Coming Up

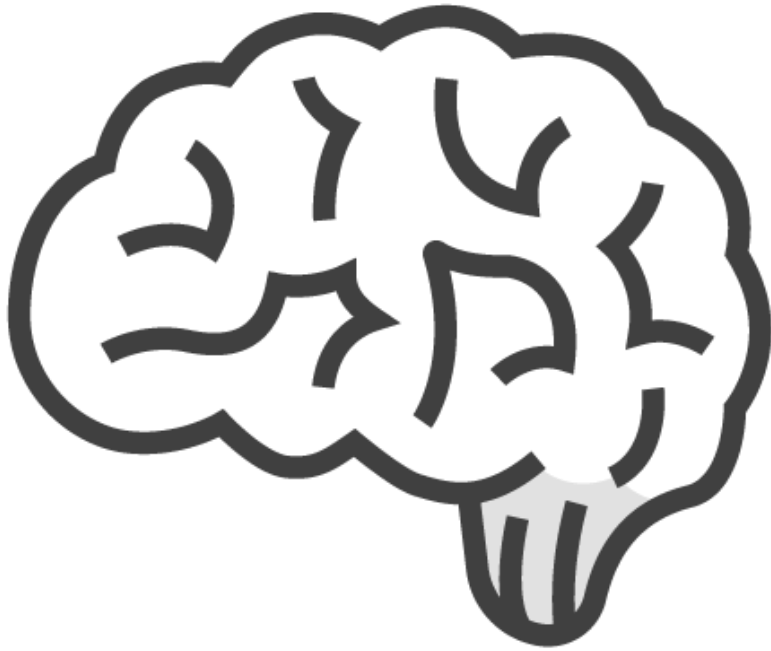


Working with caching

Concurrency in a RESTful world



Working with Caching



Each response should define itself as cacheable or not

HTTP Caching

<http://bit.ly/2hJTTxD> (RFC 2616)

<http://bit.ly/2in4uzh> (RFC 7234)

“Caching would be useless if it did not significantly improve performance. The goal of caching in HTTP/1.1 is to eliminate the need to send requests in many cases, and to eliminate the need to send full responses in many other cases.”

HTTP standard



The Purpose of Caching



Eliminate the number of requests

Reduces network-roundtrips
Expiration mechanism

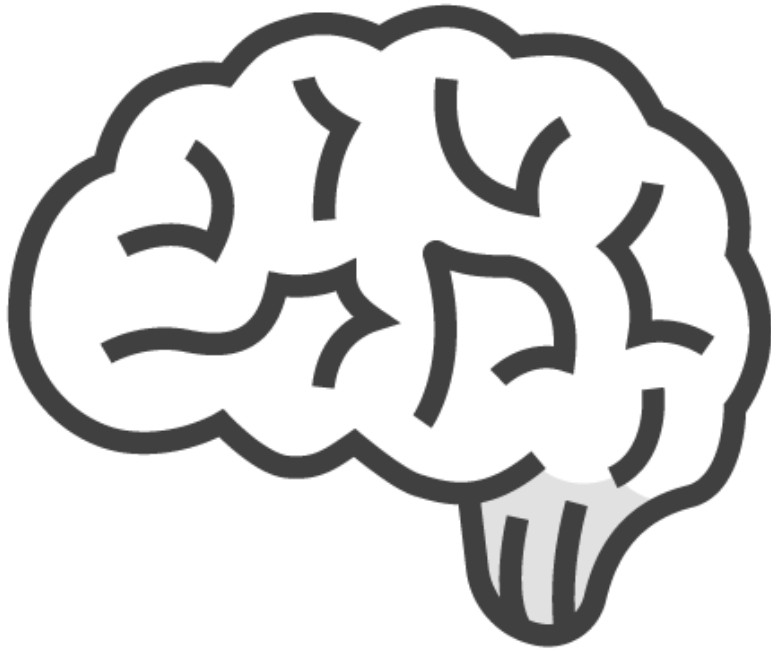


Eliminate the need to send full responses

Reduces network bandwidth
Validation mechanism



Working with Caching



The cache is a separate component

Accepts requests from consumer to the API

Receives responses from the API and stores them if they are deemed cacheable


It's the middle-man of request-response communication

Cache Types




Client Cache

Lives on the client
Private cache




Gateway Cache

Lives on the server
Shared cache



Proxy Cache

Lives on the network
Shared cache



Expiration Model

Allows the server to state how long a response is considered fresh

Expiration Model

Expires header

Expires: Sat, 14 Jan 2017 15:23:40 GMT

Clocks must be synchronised

Offers little control

Cache-Control header

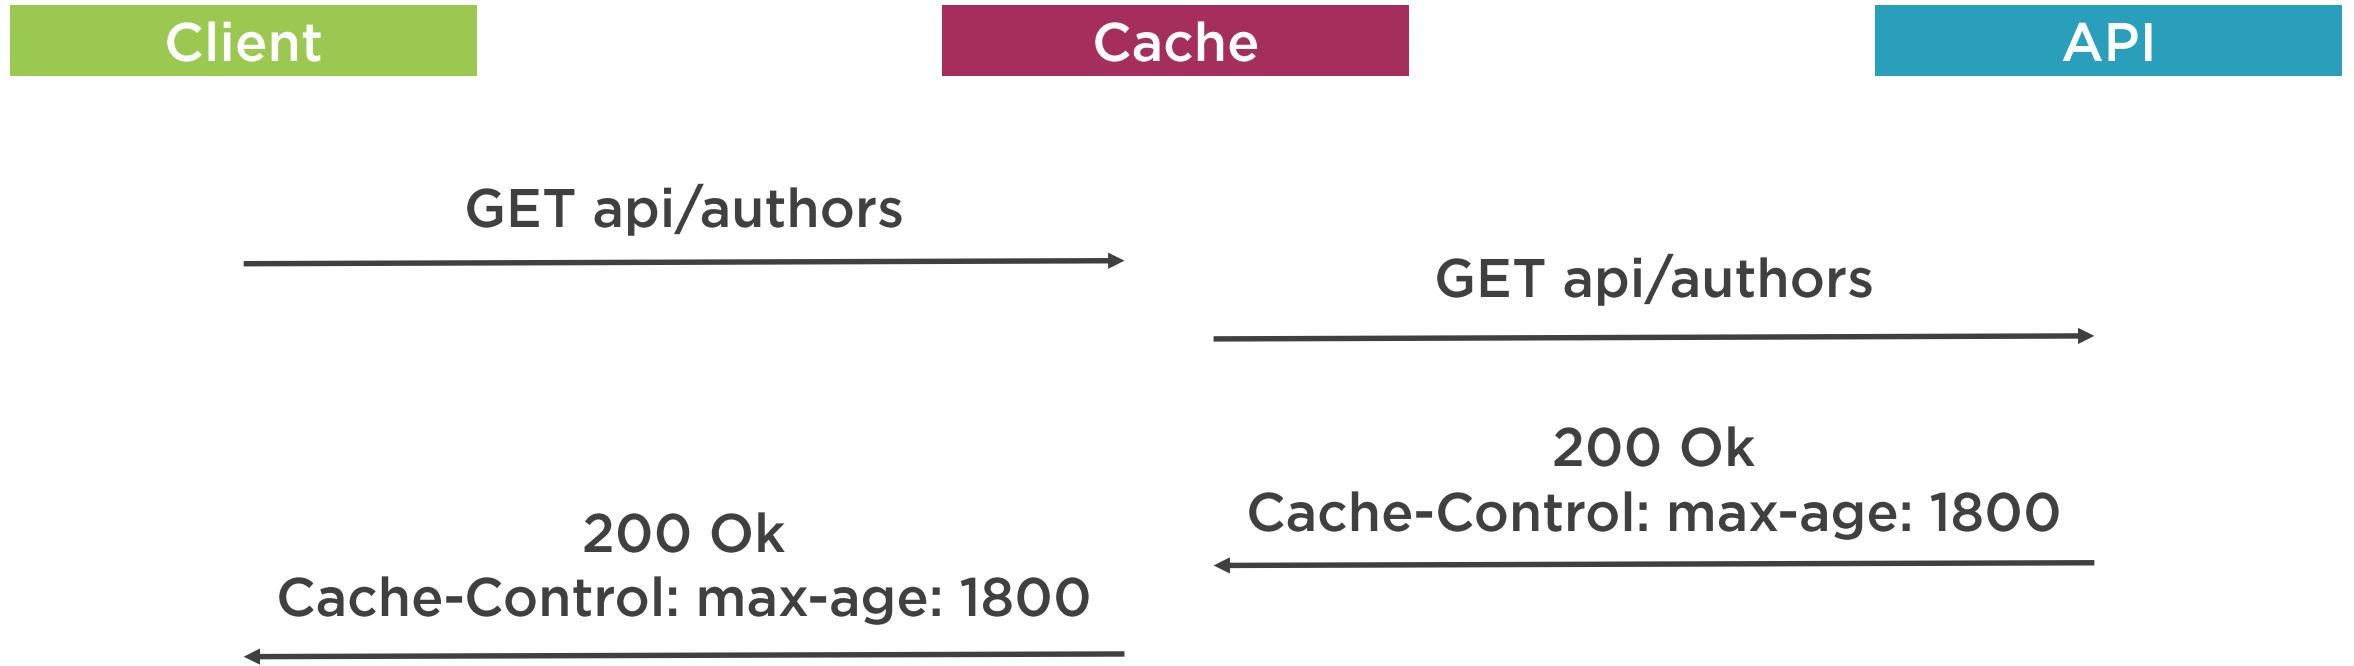
Cache-Control: public, max-age=60

Preferred header for expiration

Directives: <http://bit.ly/1Ups120>



Expiration Model



Expiration Model

Client

Cache

API

GET api/authors



```
graph LR; Client[Client] -- "GET api/authors" --> Cache[Cache]; Cache -- "200 Ok<br/>Age: 600<br/>Cache-Control: max-age: 1800" --> Client;
```

200 Ok
Age: 600

Cache-Control: max-age: 1800



Expiration Model

Client

Cache

API

GET api/authors



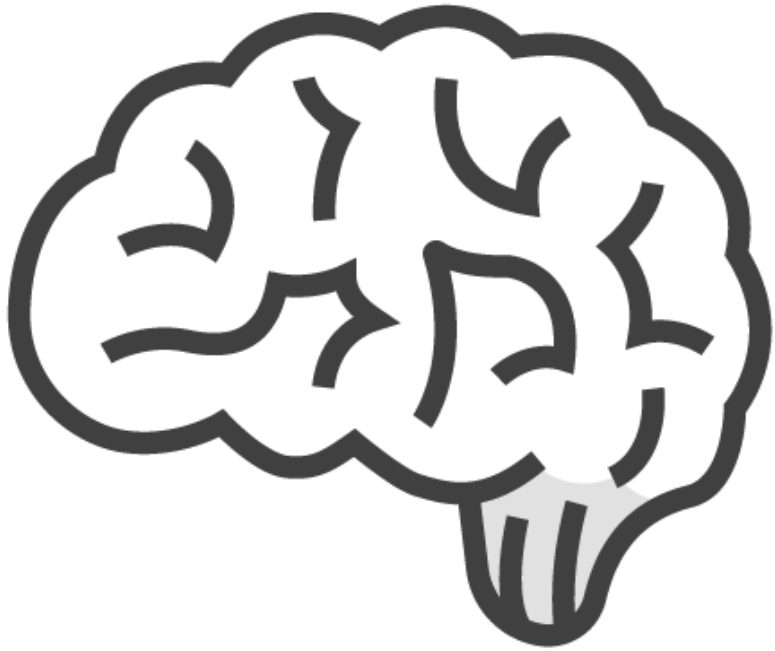
```
graph LR; Client[Client] -- "GET api/authors" --> Cache[Cache]; Cache -- "200 Ok<br/>Age: 1200<br/>Cache-Control: max-age: 1800" --> Client;
```

200 Ok
Age: 1200

Cache-Control: max-age: 1800



Expiration Model and Cache Types




Private cache

Reduces bandwidth requirements
Less requests from cache to API

Shared (public) cache

Doesn't save bandwidth between cache and API
Drastically lowers requests to the API



Validation Model

Used to validate the freshness of a response that's been cached



Validators

Strong validators

Change if the body or headers of a response change

ETag (Entity Tag) response header
ETag: "123456789"

Can be used in any context (equality is guaranteed)

Weak validators

Don't always change when the response changes (eg: only on significant changes)

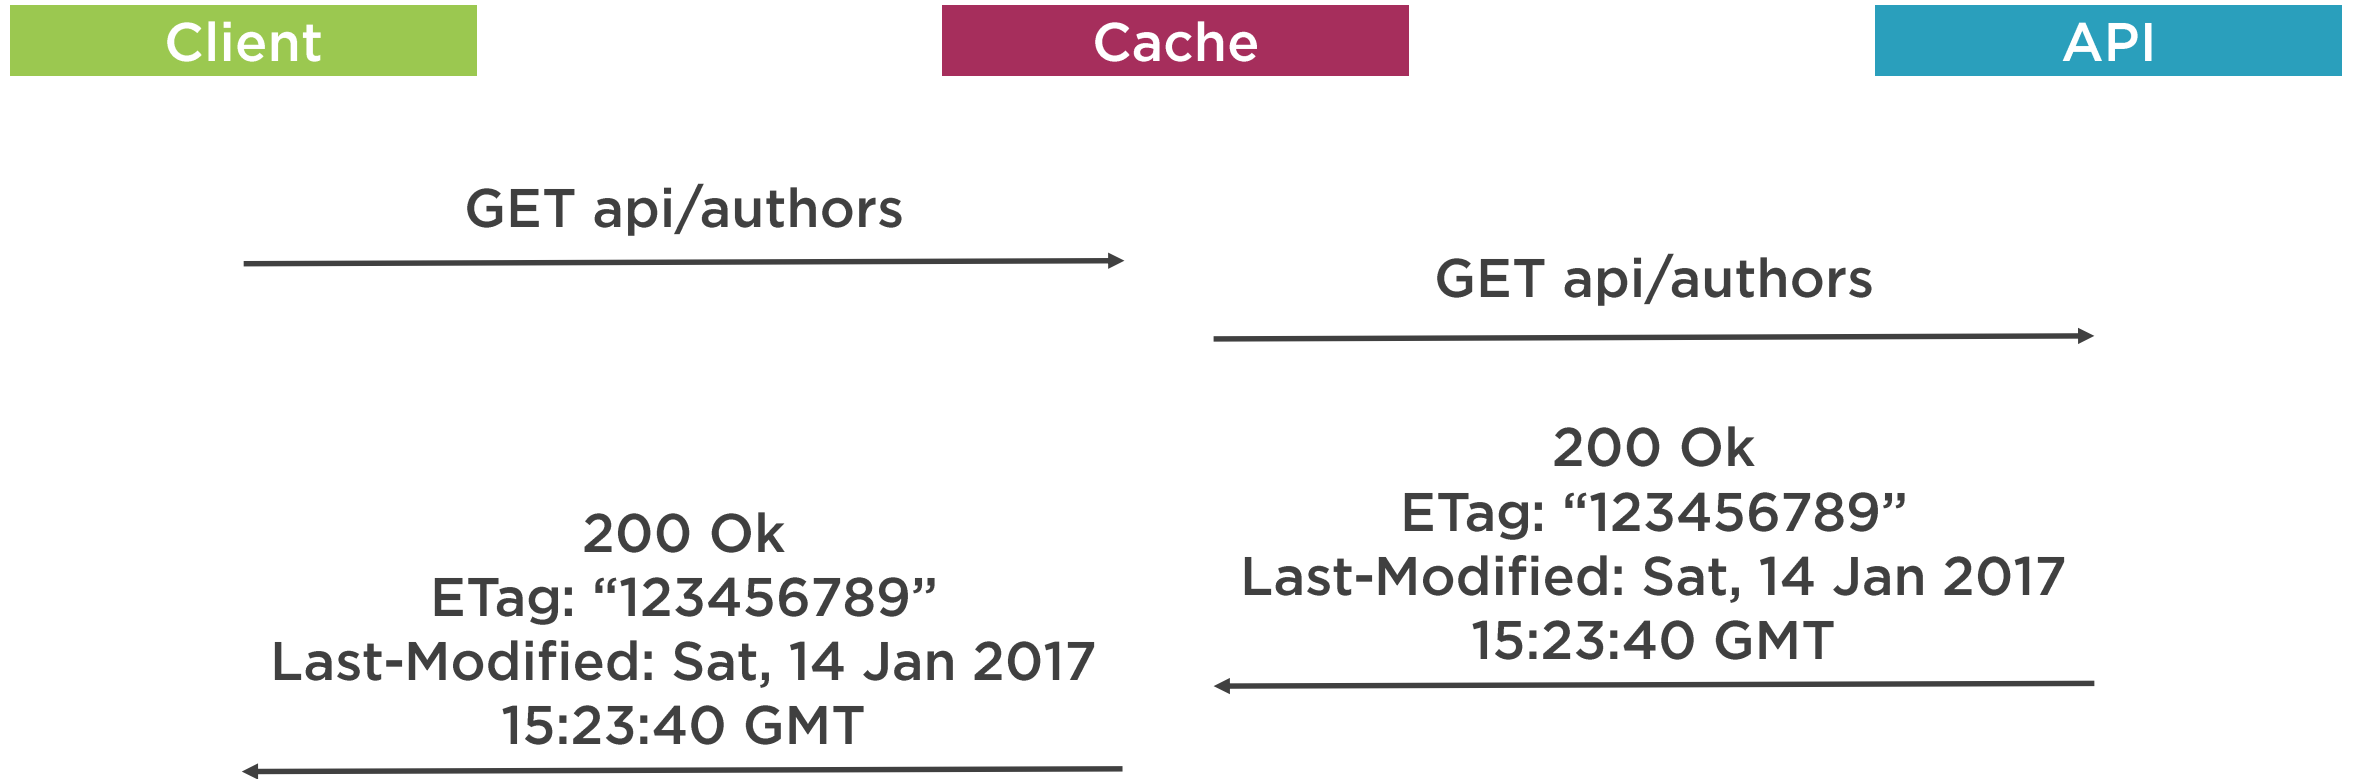
Last-Modified: Sat, 14 Jan 2017 15:23:40 GMT

ETag: "w/123456789"

Equivalence, but not equality



Validation Model



Validation Model

Client

Cache

API

GET api/authors

GET api/authors
If-None-Match: "123456789"
If-Modified-Since: Sat, 14 Jan
2017 15:23:40 GMT

304 Not Modified

200 Ok
ETag: "123456789"
Last-Modified: Sat, 14 Jan 2017
15:33:40 GMT



Validation Model

Client

Cache

API

GET api/authors

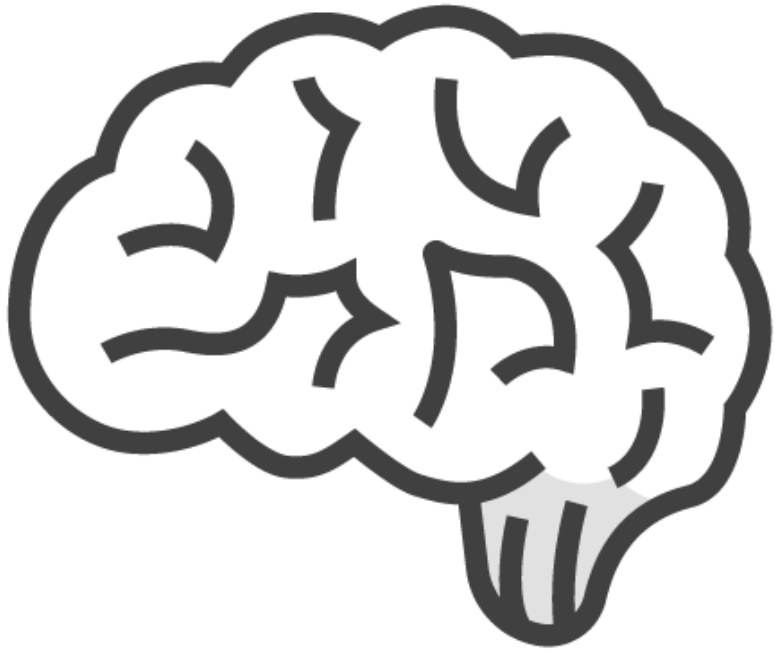
GET api/authors
If-None-Match: "123456789"
If-Modified-Since: Sat, 14 Jan
2017 15:33:40 GMT

304 Not Modified

200 Ok
ETag: "123456789"
Last-Modified: Sat, 14 Jan 2017
15:43:40 GMT



Validation Model and Cache Types



Private cache

Reduces bandwidth requirements

Shared (public) cache

Saves bandwidth between cache and API



Expiration and Validation Combined

Private cache

As long as the response hasn't expired (isn't stale), that response can be returned from the cache

Reduces communication with the API (including response generation), reduces bandwidth requirements

If it has expired, the API is hit

Bandwidth usage and response generation is potentially reduced even more

Shared (public) cache

As long as the response hasn't expired (isn't stale), that response can be returned from the cache

Reduces bandwidth requirements between cache and API, dramatically reduces request to the API

If it has expired, the API is hit

Bandwidth usage between cache and API and response generation is potentially reduced





The Holy Grail of Caching

Combine private and shared caches



Cache-Control Directives

Response

Freshness

max-age, s-maxage

Cache type

public, private

Validation

no-cache, must-revalidate,
proxy-revalidate

Other

no-store, no-transform

Request

Freshness

max-age, min-fresh, max-stale

Validation

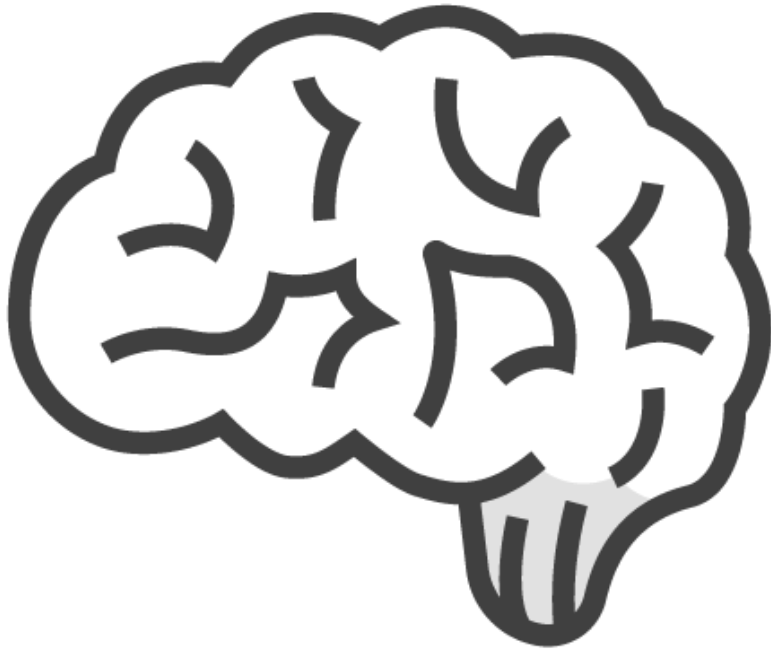
no-cache

Other

no-store, no-transform, only-if-cached



Supporting Cache Headers



In ASP.NET Core: [ResponseCache]

Not sufficient for our purposes

In ASP.NET: CacheCow.Server

Not available for ASP.NET Core at the moment

Marvin.Cache.Headers

- NuGet: <http://bit.ly/2knD67M>


Demo



Supporting HTTP Cache Headers



Cache Stores




Private caches

- angular-http-etag (<http://bit.ly/2jqm6hh>)
- Marvin.HttpCache (<http://bit.ly/2i6AuJE>)
 - PCL, no .NET Core

Private and shared caches

- CacheCow.Client (<http://bit.ly/2iqWIUs>)
 - ASP.NET Web API, full .NET framework

Cache Stores



Shared caches (.NET Core)

- ASP.NET Core ResponseCaching middleware (<http://bit.ly/2jq4UbM>)
- Use from v2.0 onwards

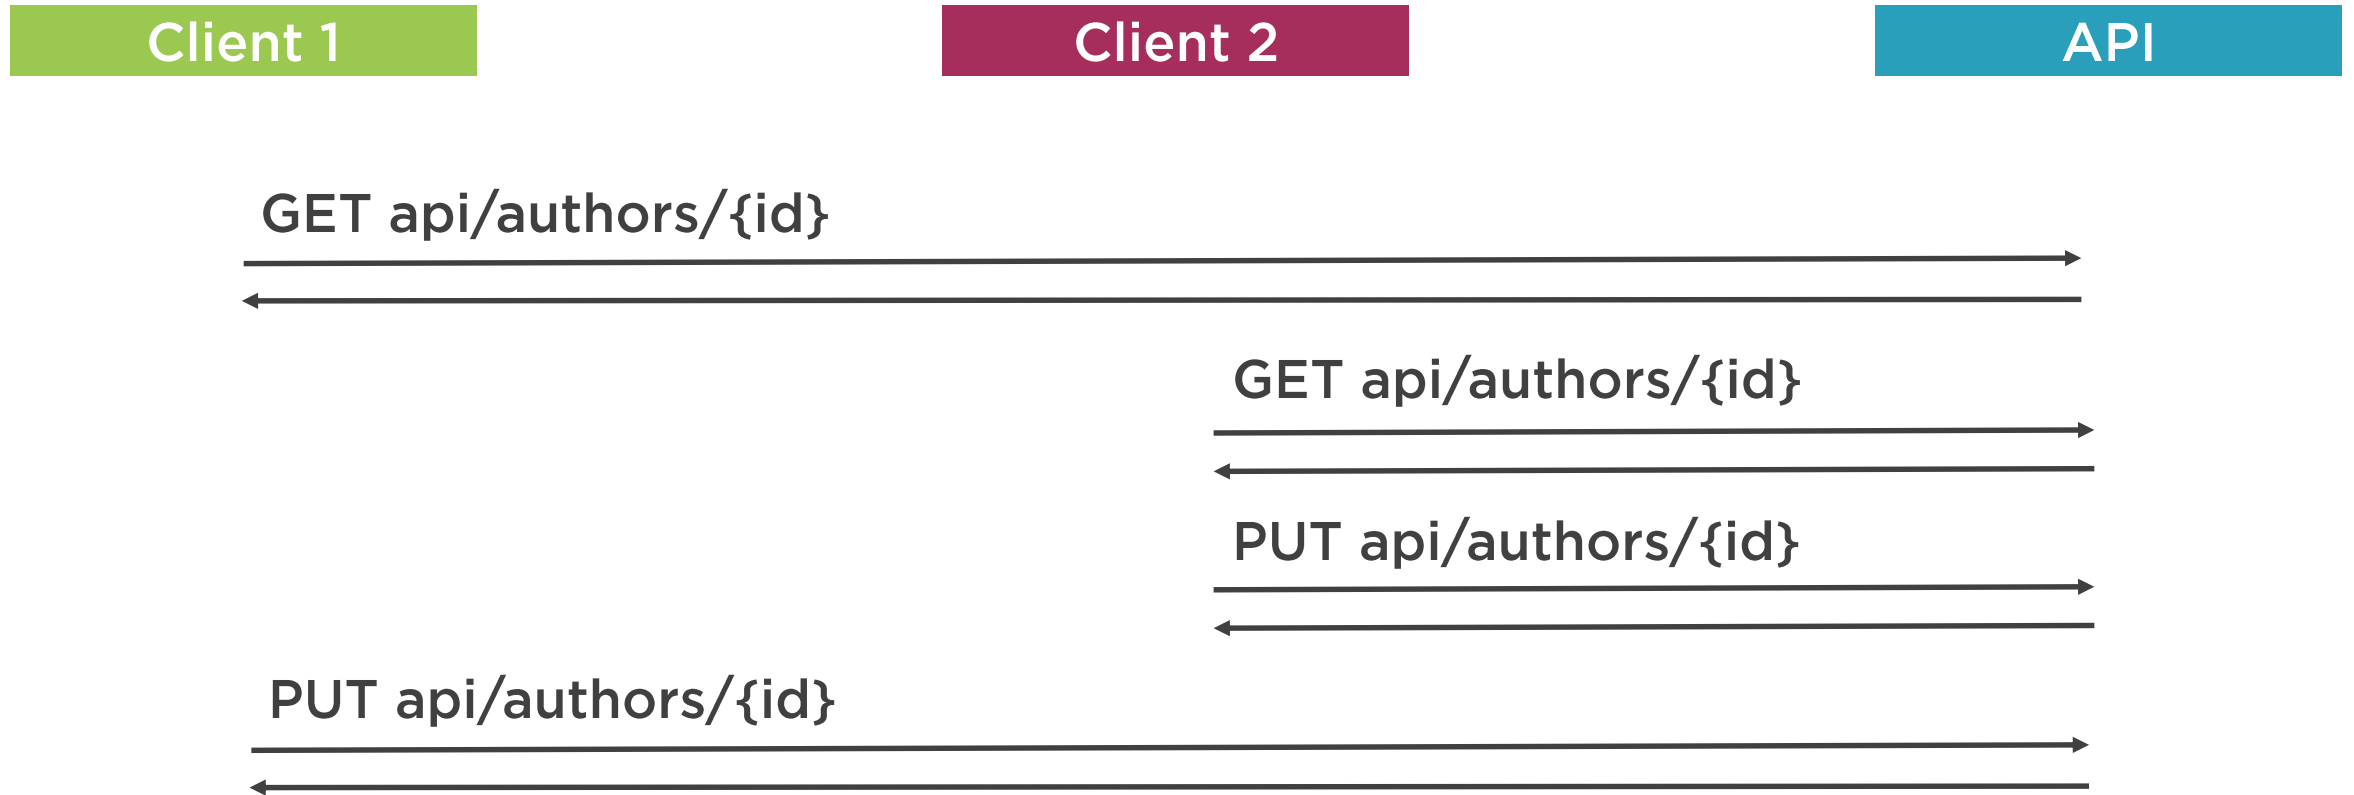
Demo



Adding a Cache Store



Dealing with Concurrency in a RESTful World



Concurrency Strategies

Pessimistic concurrency

Resource is locked

While it's locked, it cannot be modified
by another client

This is not possible in REST

Optimistic concurrency

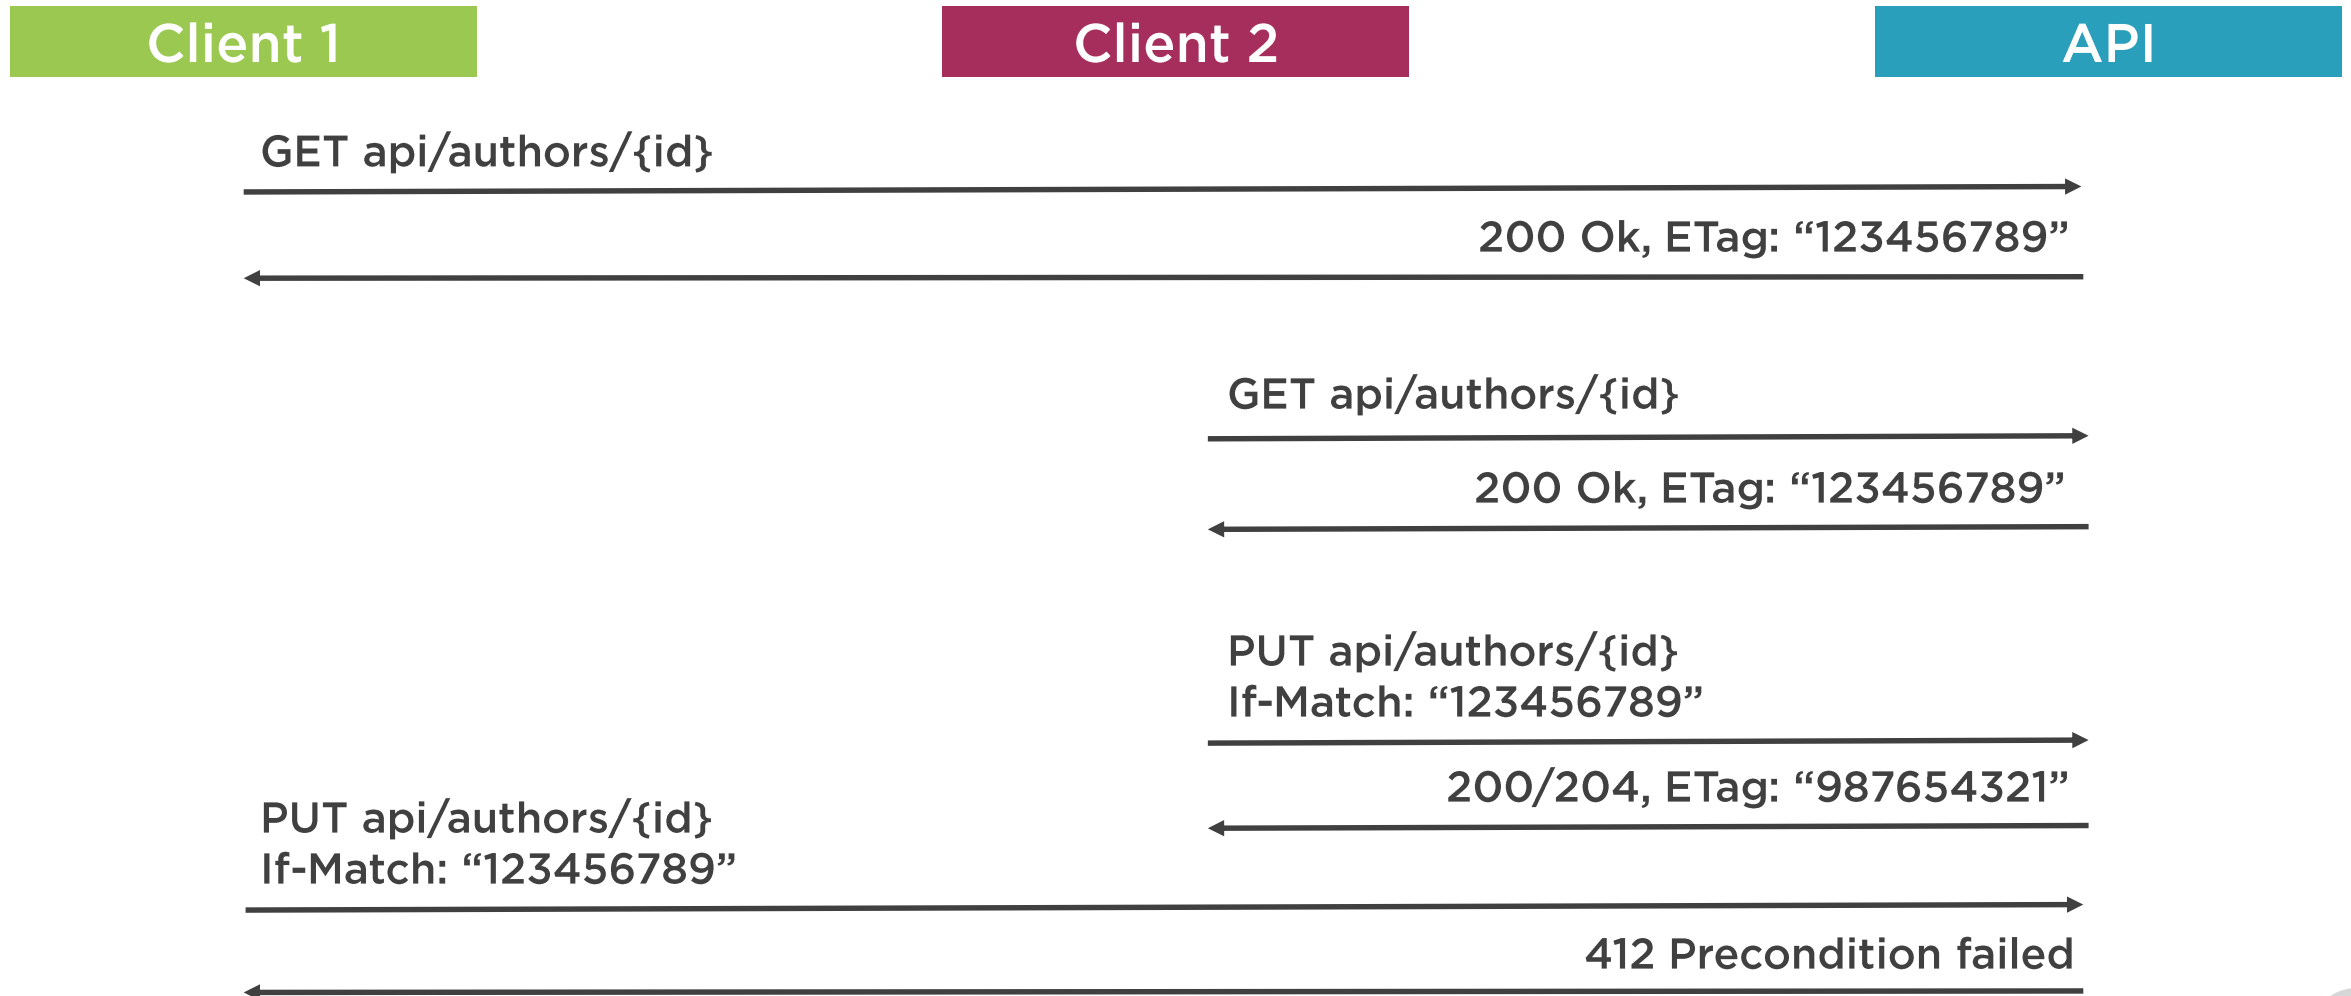
Token is returned together with the
resource

The update can happen as long as the
token is still valid

ETags are used as validation tokens



Dealing with Concurrency in a RESTful World



Demo



Dealing with Concurrency



Summary



Each response must state whether or not it can be cached

Caching: expiration model

Allows the server to state how long a response is considered fresh

Cache-Control header

Summary



Caching: validation model

Used to validate the freshness of a response that's been cached

ETag, Last-Modified headers

Use ETags for optimistic concurrency